

PySpark

Ming-Hwa Wang, Ph.D.
COEN 242 Introduction to Big Data
Department of Computer Engineering
Santa Clara University

<https://github.com/PacktPublishing/Learning-PySpark/>

Load RDDs

- Read JSON file:
 - Read JSON file from sqlContext:

```
df = sqlContext.read \
    .format('json').load('py/test/sql/people.json')
```
 - Read JSON file from spark.read.format('json').load(f) or spark.read.json(f):

```
df = spark.read.format('json').load('py/test/sql/people.json')
df = spark.read.json('py/test/sql/people.json')
```
- Create RDDs:
 - .parallelize(...) a collection (tuple, list, dict, etc.):

```
data = sc.parallelize(
    [('Amber', 22), ('Alfred', 23), ('Ske', 4), ('Albert', 12),
    ('Amber', 9)]
).collect()
```
 - reference a file (or zipped/unzipped files) with number of partitions:

```
data_from_file = sc.\
    textFile(
        '/Users/drabast/Documents/PySpark_Data/VS14MORT.txt.gz', 4
    ).take(1)
```
- Extract information with Python method (can slow down your application):

```
def extractInformation(row):
    import re
    import numpy as np
    selected_indices = [
        2,4,5,6,7,9,10,11,12,13,14,15,16,17,18,...
        77,78,79,81,82,83,84,85,87,89
    ]
    record_split = re\
    .compile(
        r'([\s]{19})([0-9]{1})([\s]{40})...
        ([\s]{33})([0-9\s]{3})([0-9\s]{1})([0-9\s]{1})'
    )
    try:
        rs = np.array(record_split.split(row))[selected_indices]
    except:
        rs = np.array(['-99'] * len(selected_indices))
    return rs
```
- Filter out the malformed records:

```
data_from_file_conv = data_from_file.map(extractInformation)
```

Transformations

- The .map() transformation applies to each element of the dataset:

```
data_2014 = data_from_file_conv.map(lambda row: int(row[16]))
```
- The .filter() transformation select elements that fit specified criteria:

```
data_filtered = data_from_file_conv.filter(
    lambda row: row[16] == '2014' and row[21] == '0')
```
- The .flatMap() transformation returns a flattened result instead of a list:

```
data_2014_flat = data_from_file_conv.flatMap(
    lambda row: (row[16], int(row[16]) + 1))
```
- The .distinct() transformation returns a list of distinct values in a column (an expensive transformation):

```
distinct_gender = data_from_file_conv.map(
    lambda row: row[5]).distinct()
```
- The .sample() transformation returns a randomized sample from the dataset with replacement:

```
fraction = 0.1
seed = 666
data_sample = data_from_file_conv.sample(False, fraction, seed)
```
- The .join() transformation returns records when the two RDDs match (an expensive transformation):

```
rdd1 = sc.parallelize([('a', 1), ('b', 4), ('c',10)])
rdd2 = sc.parallelize([('a', 4), ('a', 1), ('b', '6'), ('d', 15)])
rdd3 = rdd1.join(rdd2)
produce the following:
[ ('b', (4, '6')), ('a', (1, 4)), ('a', (1, 1))]
```
- The .leftOuterJoin() transformation returns records from the left RDD with records from the right one appended in places where the two RDDs match (an expensive transformation):

```
rdd1 = sc.parallelize([('a', 1), ('b', 4), ('c',10)])
rdd2 = sc.parallelize([('a', 4), ('a', 1), ('b', '6'), ('d', 15)])
rdd3 = rdd1.leftOuterJoin(rdd2)
produce the following:
[(('c', (10, None)), ('b', (4, '6')), ('a', (1, 4)), ('a', (1, 1)))]
```
- The .intersection() transformation returns records that are equal in both RDDs:

```
rdd5 = rdd1.intersection(rdd2)
rdd5.collect()
produce the following:
[('a', 1)]
```
- The .repartition() transformation changes the number of partitions (an expensive transformation):

```
rdd1 = rdd1.repartition(4)
len(rdd1.glom().collect())
```

The `.glom()` method produces a list where each element is another list of all elements of the dataset present in a specified partition.

Actions

Actions execute the scheduled task on the dataset. This might contain no transformations.

- The `.take()` method returns the n top rows from a single data partition:
`data_first = data_from_file_conv.take(1)`
- The `.takeSample()` method returns the n random sample rows from a single data partition:
`replacement = False`
`n = 1`
`seed = 667`
`data_take_sampled = data_from_file_conv.takeSample(replacement, n, seed)`
- The `.collect()` method returns all the elements of the RDD to the driver:
- The `.reduce()` method reduces the elements of an RDD using a specified method:
`rdd1.map(lambda row: row[1]).reduce(lambda x, y: x + y)`
The functions passed as a reducer need to be associative and commutative.
- The `.reduceByKey()` method performs a reduction on a key-by-key basis:
`data_key = sc.parallelize([('a', 4), ('b', 3), ('c', 2), ('a', 8), ('d', 2), ('b', 1), ('d', 3)], 4)`
`data_key.reduceByKey(lambda x, y: x + y).collect()`
- The `.count()` method counts the number of elements in the RDD.
`data_reduce.count()`
- The `.countByKey()` method gets the counts of distinct keys:
`data_key.countByKey().items()`
- The `.saveAsTextFile()` method saves to text files, each partition to a separate file:
`data_key.saveAsTextFile('/Users/drabast/Documents/PySpark_Data/data_key.txt')`
 - To read it back, parse it back as all the rows are treated as strings:

```
def parseInput(row):  
    import re  
    pattern = re.compile(r'\\(\\'([a-z])\\', ([0-9])\\')  
    row_split = pattern.split(row)  
    return (row_split[1], int(row_split[2]))  
data_key_reread = sc.textFile(  
    '/Users/drabast/Documents/PySpark_Data/data_key.txt') \br/>    .map(parseInput)  
data_key_reread.collect()
```
- The `.foreach()` method applies the same function to each element of the RDD:

```
def f(x):  
    print(x)  
data_key.foreach(f)
```

DataFrames

- Generating JSON data:

```
stringJSONRDD = sc.parallelize(  
    ""{  
        "id": "123",  
        "name": "Katie",  
        "age": 19,  
        "eyeColor": "brown"  
    }""  
    ""{  
        "id": "234",  
        "name": "Michael",  
        "age": 22,  
        "eyeColor": "green"  
    }""  
    ""{  
        "id": "345",  
        "name": "Simone",  
        "age": 23,  
        "eyeColor": "blue"  
    }""  
))
```
- Creating a DataFrame:
`swimmersJSON = spark.read.json(stringJSONRDD)`
- Creating a temporary table:
`swimmersJSON.createOrReplaceTempView("swimmersJSON")`
- DataFrame API query using the `show(n=10)` method:
`swimmersJSON.show()`
- SQL query
`spark.sql("select * from swimmersJSON").collect()`
 - Databricks uses the `%sql` command and run your SQL statement directly within a notebook cell.
- Interoperating with RDDs: programmatically specifying the schema:

```
# Import types  
from pyspark.sql.types import *  
# Generate comma delimited data  
stringCSRDD = sc.parallelize([  
    (123, 'Katie', 19, 'brown'),  
    (234, 'Michael', 22, 'green'),  
    (345, 'Simone', 23, 'blue')  
])  
# Specify schema  
schema = StructType([  
    StructField("id", LongType(), True),  
    StructField("name", StringType(), True),  
    StructField("age", LongType(), True),
```

```

    StructField("eyeColor", StringType(), True)
])
# Apply the schema to the RDD and Create DataFrame
swimmers = spark.createDataFrame(stringCSVRDD, schema)
# Creates a temporary view using the DataFrame
swimmers.createOrReplaceTempView("swimmers")
# find schema from DataFrame
swimmers.printSchema()

```

- Querying with the DataFrame


```

# get the number of rows
swimmers.count()
# Running filter statements. Get the id, age where age = 22
swimmers.select("id", "age").filter("age = 22").show()
# Another way to write the above query is below
swimmers.select(swimmers.id, swimmers.age) \
    .filter(swimmers.age == 22).show()
# Get the name, eyeColor where eyeColor like 'b%'
swimmers.select("name", "eyeColor") \
    .filter("eyeColor like 'b%'").show()

```
- Querying with SQL


```

#number of rows
spark.sql("select count(1) from swimmers").show()
# running filter statements using the where Clauses
# Get the id, age where age = 22 in SQL
spark.sql("select id, age from swimmers where age = 22").show()
spark.sql(
    "select name, eyeColor from swimmers where eyeColor like 'b%'")
    .show()

```
- DataFrame scenario example – on-time flight performance.
 - Preparing the source datasets


```

# Set File Paths
flightPerfFilePath =
    "/databricks-datasets/flights/departuredelays.csv"
airportsFilePath =
    "/databricks-datasets/flights/airport-codes-na.txt"
# Obtain Airports dataset
airports = spark.read.csv(
    airportsFilePath, header='true', inferSchema='true',
    sep='\t'
)
airports.createOrReplaceTempView("airports")
# Obtain Departure Delays dataset
# csv file can be common-delimited (default) or tab-delimited
flightPerf = spark.read.csv(flightPerfFilePath, header='true')
flightPerf.createOrReplaceTempView("FlightPerformance")
# Cache the Departure Delays dataset
flightPerf.cache()

```
 - Joining flight performance and airports

```

# Query Sum of Flight Delays by City and Origin Code
# (for Washington State)
spark.sql("""
    select a.City, f.origin, sum(f.delay) as Delays
    from FlightPerformance f
    join airports a
    on a.IATA = f.origin
    where a.State = 'WA'
    group by a.City, f.origin
    order by sum(f.delay) desc""")
    ).show()

```

- Databricks notebook can use the %sql function to execute SQL statements within that notebook cell to get the same result as previous query but easier to read:


```

%sql
-- Query Sum of Flight Delays by City and Origin Code (for
-- Washington State)
select a.City, f.origin, sum(f.delay) as Delays
from FlightPerformance f
join airports a
on a.IATA = f.origin
where a.State = 'WA'
group by a.City, f.origin
order by sum(f.delay) desc

```
- Visualizing our flight-performance data

Getting familiar with your data

We can build a model without knowing data by taking longer time with suboptimal results. Thus, any serious data scientist or data modeler will become acquainted with the dataset before starting any modeling.

- Descriptive statistics tell you the basic information about your dataset: mean, standard deviation, min and max.


```

import pyspark.sql.types as typ
fraud = sc.textFile('ccFraud.csv.gz')
header = fraud.first()
fraud = fraud.filter(lambda row: row != header) \
    .map(lambda row: [int(elem) for elem in row.split(',')])
# create the schema for our DataFrame:
fields = [*[
    typ.StructField(h[1:-1], typ.IntegerType(), True) \
    for h in header.split(',')
]]
schema = typ.StructType(fields)
# create our DataFrame:
fraud_df = spark.createDataFrame(fraud, schema)
# get the schema of our DataFrame:
fraud_df.printSchema()

```

```
# for categorical features, count the frequencies of their values
fraud_df.groupby('gender').count().show()
# for the truly numerical features, use the .describe() method
numerical = ['balance', 'numTrans', 'numIntlTrans']
desc = fraud_df.describe(numerical)
desc.show()
# check the skewness
fraud_df.agg({'balance': 'skewness'}).show()
A list of aggregation functions (the names are fairly self-explanatory) includes:
avg(), count(), countDistinct(), first(), kurtosis(), max(), mean(), min(),
skewness(), stddev(), stddev_pop(), stddev_samp(), sum(), sumDistinct(),
var_pop(), var_samp() and variance().
```

- Correlations

```
# calculate pairwise correlations:
fraud_df.corr('balance', 'numTrans')
# create a correlations matrix:
n_numerical = len(numerical)
corr = []
for i in range(0, n_numerical):
    temp = [None] * i
    for j in range(i+1, n_numerical):
        temp.append(fraud_df.corr(numerical[i], numerical[j]))
    corr.append(temp)
```

- Visualization using visualization packages: matplotlib and Bokeh (preinstalled with Anaconda).

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('ggplot')
import bokeh.charts as chrt
from bokeh.io import output_notebook
output_notebook()
```

- Histograms: visualize distribution of your features

```
# aggregate the data first:
hists = fraud_df.select('balance').rdd.flatMap(
    lambda row: row
).histogram(20)
# plot the histogram using matplotlib:
data = {
    'bins': hists[0][:-1], 'freq': hists[1]
}
plt.bar(data['bins'], data['freq'], width=2000)
plt.title('Histogram of \'balance\'')
# histogram created with Bokeh (which uses interactive D3.js):
b_hist = chrt.Bar(
    data, values='freq', label='bins',
    title='Histogram of \'balance\'')
)
```

```
chrt.show(b_hist)
# small data, use matplotlib's .hist() or Bokeh's .Histogram()
data_driver = {
    'obs': fraud_df.select('balance').rdd.flatMap(
        lambda row: row
    ).collect()
}
plt.hist(data_driver['obs'], bins=20)
plt.title('Histogram of \'balance\' using .hist()')
b_hist_driver = chrt.Histogram(
    data_driver, values='obs',
    title='Histogram of \'balance\' using .Histogram()',
    bins=20
)
chrt.show(b_hist_driver)
```

- Interactions between features: scatter charts for up to 3 variables, 3D visualizations with temporal data, and sample huge dataset.

```
# samples at a predefined sampling fraction
data_sample = fraud_df.sampleBy(
    'gender', {1: 0.0002, 2: 0.0002}
).select(numerical)
# put multiple 2D charts in one:
data_multi = dict([
    (elem, data_sample.select(elem).rdd \
     .flatMap(lambda row: row).collect()) for elem in numerical
])
sctr = chrt.Scatter(data_multi, x='balance', y='numTrans')
chrt.show(sctr)
```

Prepare Data for Modeling

Checking for duplicates, missing observations, and outliers

- Duplicates removal with .distinct() method:

```
# example dataset:
df = spark.createDataFrame(
    [
        (1, 144.5, 5.9, 33, 'M'),
        (2, 167.2, 5.4, 45, 'M'),
        (3, 124.1, 5.2, 23, 'F'),
        (4, 144.5, 5.9, 33, 'M'),
        (5, 133.2, 5.7, 54, 'F'),
        (3, 124.1, 5.2, 23, 'F'),
        (5, 129.2, 5.3, 42, 'M'),
    ],
    ['id', 'weight', 'height', 'age', 'gender']
)
print('Count of rows: {0}'.format(df.count()))
print('Count of distinct rows: {0}'.format(df.distinct().count()))
# if these two numbers differ, drop these rows
```

```

df = df.dropDuplicates()
# check whether any duplicates in the data irrespective of ID
print('Count of ids: {}'.format(df.count()))
print('Count of distinct ids: {}'.format(df.select(
    [c for c in df.columns if c != 'id']
).distinct().count()))
# or use the .dropDuplicates(),
df = df.dropDuplicates(subset=[
    c for c in df.columns if c != 'id'
])
# to calculate the total and distinct number of IDs in one step
import pyspark.sql.functions as fn
df.agg(
    fn.count('id').alias('count'),
    fn.countDistinct('id').alias('distinct')
).show()
# give each row a unique ID:
df.withColumn('new_id',
    fn.monotonically_increasing_id()).show()

```

- Missing observations


```

# example dataset:
df_miss = spark.createDataFrame(
    [
        (1, 143.5, 5.6, 28, 'M', 100000),
        (2, 167.2, 5.4, 45, 'M', None),
        (3, None, 5.2, None, None, None),
        (4, 144.5, 5.9, 33, 'M', None),
        (5, 133.2, 5.7, 54, 'F', None),
        (6, 124.1, 5.2, None, 'F', None),
        (7, 129.2, 5.3, 42, 'M', 76000),
    ],
    ['id', 'weight', 'height', 'age', 'gender', 'income']
)
# find the number of missing observations
df_miss.rdd.map(
    lambda row: (row['id'], sum([c == None for c in row]))
).collect()
# produces [(1,0), (2,1), (3,4), (4,1), (5,1), (6,2), (7,0)]
df_miss.where('id == 3').show()
# percentage of missing, the * argument in .count() counts all rows
df_miss.agg(*[
    (1 - (fn.count(c) / fn.count('*))).alias(c + '_missing') \
    for c in df_miss.columns
]).show()
# drop the 'income' feature, as most of its values are missing.
df_miss_no_income = df_miss.select([
    c for c in df_miss.columns if c != 'income'
])

```

```

df_miss_no_income.dropna(thresh=3).show()
# If you want to impute a mean, median, or other calculated value,
# you need to first calculate the value, create a dictionary with
# such values, and then pass it to the .fillna() method.
# records parameter in .to_dict() instructs to create a dictionary
means = df_miss_no_income.agg(
    *[fn.mean(c).alias(c)
      for c in df_miss_no_income.columns if c != 'gender']
).toPandas().to_dict('records')[0]
means['gender'] = 'missing'
df_miss_no_income.fillna(means).show()

```

- Outliers: observations that deviate significantly from the distribution of the rest samples. There are no outliers if all the values are roughly within the Q1–1.5IQR and Q3+1.5IQR range, where IQR is the interquartile range; the IQR is defined as a difference between the upper- and lower-quartiles, that is, the 75th percentile (the Q3) and 25th percentile (the Q1), respectively.


```

# example dataset:
df_outliers = spark.createDataFrame(
    [
        (1, 143.5, 5.3, 28),
        (2, 154.2, 5.5, 45),
        (3, 342.3, 5.1, 99),
        (4, 144.5, 5.5, 33),
        (5, 133.2, 5.4, 54),
        (6, 124.1, 5.1, 21),
        (7, 129.2, 5.3, 42),
    ],
    ['id', 'weight', 'height', 'age']
)
# Calculate the lower and upper cut off points for each feature.
# The first parameter specified is the name of the column, the
# second parameter can be either a number between 0 or 1 (where
# 0.5 means to calculated median) or a list, and the third parameter
# specifies the acceptable level of an error for each metric.
cols = ['weight', 'height', 'age']
bounds = {}
for col in cols:
    quantiles = df_outliers.approxQuantile(col, [0.25, 0.75], 0.05)
    IQR = quantiles[1] - quantiles[0]
    bounds[col] = [
        quantiles[0] - 1.5 * IQR, quantiles[1] + 1.5 * IQR
    ]
# flag outliers:
outliers = df_outliers.select(
    *['id'] + [
        (
            (df_outliers[c] < bounds[c][0]) |
            (df_outliers[c] > bounds[c][1])

```

```

    ).alias(c + '_o') for c in cols
]
)
outliers.show()
# lists the values significantly differing from the rest of the
# distribution:
df_outliers = df_outliers.join(outliers, on='id')
df_outliers.filter('weight_o').select('id', 'weight').show()
df_outliers.filter('age_o').select('id', 'age').show()

```

MLlib on RDDs

For example, to predict whether the 'INFANT_ALIVE_AT_REPORT' is either 1 or 0.

- Loading and transforming the data

```

import pyspark.sql.types as typ
labels = [
    ('INFANT_ALIVE_AT_REPORT', typ.StringType()),
    ('BIRTH_YEAR', typ.IntegerType()),
    ('BIRTH_MONTH', typ.IntegerType()),
    ('BIRTH_PLACE', typ.StringType()),
    ('MOTHER_AGE_YEARS', typ.IntegerType()),
    ('MOTHER_RACE_6CODE', typ.StringType()),
    ('MOTHER_EDUCATION', typ.StringType()),
    ('FATHER_COMBINED_AGE', typ.IntegerType()),
    ('FATHER_EDUCATION', typ.StringType()),
    ('MONTH_PRECARE_RECODE', typ.StringType()),
    ...
    ('INFANT_BREASTFED', typ.StringType())
]
schema = typ.StructType([
    typ.StructField(e[0], e[1], False) for e in labels
])
births = spark.read.csv(
    'births_train.csv.gz', header=True, schema=schema
)
# specify recode dictionary:
recode_dictionary = {
    'YNU': {
        'Y': 1,
        'N': 0,
        'U': 0
    }
}

```

- Manual feature selection

```

selected_features = [
    'INFANT_ALIVE_AT_REPORT', 'BIRTH_PLACE', 'MOTHER_AGE_YEARS',
    'FATHER_COMBINED_AGE', 'CIG_BEFORE', 'CIG_1_TRI', 'CIG_2_TRI',
    'CIG_3_TRI', 'MOTHER_HEIGHT_IN', 'MOTHER_PRE_WEIGHT',
    'MOTHER_DELIVERY_WEIGHT', 'MOTHER_WEIGHT_GAIN',

```

```

    'DIABETES_PRE', 'DIABETES_GEST', 'HYP_TENS_PRE',
    'HYP_TENS_GEST', 'PREV_BIRTH_PRETERM'
]
births_trimmed = births.select(selected_features)
# specify recoding methods:
import pyspark.sql.functions as func
def recode(col, key):
    return recode_dictionary[key][col]
def correct_cig(feats):
    return func.when(
        func.col(feats) != 99, func.col(feats)
    ).otherwise(0)
# the recode function needs to be converted to a UDF on a DataFrame
rec_integer = func.udf(recode, typ.IntegerType())
# correct the features related to the number of cigarettes smoked:
births_transformed = births_trimmed \
    .withColumn('CIG_BEFORE', correct_cig('CIG_BEFORE')) \
    .withColumn('CIG_1_TRI', correct_cig('CIG_1_TRI')) \
    .withColumn('CIG_2_TRI', correct_cig('CIG_2_TRI')) \
    .withColumn('CIG_3_TRI', correct_cig('CIG_3_TRI'))
# correcting the Yes/No/Unknown features
cols = [
    (col.name, col.dataType) for col in births_trimmed.schema
]
YNU_cols = []
for i, s in enumerate(cols):
    if s[1] == typ.StringType():
        dis = births.select(s[0]).distinct().rdd \
            .map(lambda row: row[0]).collect()
        if 'Y' in dis:
            YNU_cols.append(s[0])
# DataFrames can transform the features in bulk while
# selecting features.
births.select([
    'INFANT_NICU_ADMISSION',
    rec_integer(
        'INFANT_NICU_ADMISSION', func.lit('YNU')
    ).alias('INFANT_NICU_ADMISSION_RECODE')
]).take(5)
# transform all the YNU_cols in one go
exprs_YNU = [
    rec_integer(x, func.lit('YNU')).alias(x) if x in YNU_cols
    else x
    for x in births_transformed.columns
]
births_transformed = births_transformed.select(exprs_YNU)
# check if we got it correctly:
births_transformed.select(YNU_cols[-5:]).show(5)

```


- Understanding the data


```
# calculate the descriptive statistics of the numeric features
import pyspark.mllib.stat as st
import numpy as np
numeric_cols = [
    'MOTHER_AGE_YEARS', 'FATHER_COMBINED_AGE',
    'CIG_BEFORE', 'CIG_1_TRI', 'CIG_2_TRI', 'CIG_3_TRI',
    'MOTHER_HEIGHT_IN', 'MOTHER_PRE_WEIGHT',
    'MOTHER_DELIVERY_WEIGHT', 'MOTHER_WEIGHT_GAIN'
]
numeric_rdd = births_transformed.select(numeric_cols) \
    .rdd.map(lambda row: [e for e in row])
mllib_stats = st.Statistics.colStats(numeric_rdd)
for col, m, v in zip(
    numeric_cols, mllib_stats.mean(), mllib_stats.variance()
):
    print(
        '{0}: \t{1:.2f} \t {2:.2f}'.format(col, m, np.sqrt(v))
    )
# For categorical variables, calculate the frequencies of values:
categorical_cols = [
    e for e in births_transformed.columns if e not in numeric_cols
]
categorical_rdd = births_transformed.select(categorical_cols) \
    .rdd.map(lambda row: [e for e in row])
for i, col in enumerate(categorical_cols):
    agg = categorical_rdd.groupBy(lambda row: row[i]) \
        .map(lambda row: (row[0], len(row[1])))
    print(
        col, sorted(agg.collect(), key=lambda el: el[1],
            reverse=True)
    )
```
- Correlations: identify collinear numeric features


```
corrs = st.Statistics.corr(numeric_rdd)
for i, el in enumerate(corrs > 0.5):
    correlated = [
        (numeric_cols[j], corrs[i][j]) for j, e in enumerate(el)
        if e == 1.0 and j != i
    ]
    if len(correlated) > 0:
        for e in correlated:
            print(
                '{0}-to-{1}: {2:.2f}' \
                .format(numeric_cols[i], e[0], e[1])
            )
features_to_keep = [
    'INFANT_ALIVE_AT_REPORT', 'BIRTH_PLACE', 'MOTHER_AGE_YEARS',
    'FATHER_COMBINED_AGE', 'CIG_1_TRI', 'MOTHER_HEIGHT_IN',
```

- ```
'MOTHER_PRE_WEIGHT', 'DIABETES_PRE', 'DIABETES_GEST',
'HYP_TENS_PRE', 'HYP_TENS_GEST', 'PREV_BIRTH_PRETERM'
]
births_transformed = births_transformed.select([
 e for e in features_to_keep
])
```
- Statistical testing
 

```
for categorical features, run a Chi-square test to determine
if there are significant differences.
import pyspark.mllib.linalg as ln
for cat in categorical_cols[1:]:
 agg = births_transformed.groupby('INFANT_ALIVE_AT_REPORT') \
 .pivot(cat).count()
 agg_rdd = agg.rdd.map(lambda row: (row[1:])) \
 .flatMap(lambda row: [0 if e == None else e for e in row]) \
 .collect()
 row_length = len(agg.collect()[0]) - 1
 agg = ln.Matrices.dense(row_length, 2, agg_rdd)
 test = st.Statistics.chiSqTest(agg)
 print(cat, round(test.pValue, 4))
print(ln.Matrices.dense(3,2, [1,2,3,4,5,6]))
```
  - Creating an RDD of LabeledPoints
 

```
use a hashing trick to encode the 'BIRTH_PLACE' feature:
import pyspark.mllib.feature as ft
import pyspark.mllib.regression as reg
hashing = ft.HashingTF(7)
births_hashed = births_transformed.rdd \
 .map(lambda row: [
 list(hashing.transform(row[1]).toArray())
 if col == 'BIRTH_PLACE'
 else row[i]
 for i, col in enumerate(features_to_keep)
]).map(lambda row: [
 [e] if type(e) == int else e for e in row
]).map(lambda row: [
 item for sublist in row for item in sublist
]).map(lambda row: reg.LabeledPoint(
 row[0], ln.Vectors.dense(row[1:]))
)
```
  - Splitting into training and testing
 

```
births_train, births_test = births_hashed.randomSplit([0.6, 0.4])
```
  - Predicting infant survival by logistic regression with SGD
 

```
from pyspark.mllib.classification \
import LogisticRegressionWithLBFGS
LR_Model = LogisticRegressionWithLBFGS \
 .train(births_train, iterations=10)
LR_results = (
 births_test.map(lambda row: row.label).zip(
```

```

 LR_Model.predict(births_test.map(
 lambda row: row.features)
)
).map(
 lambda row: (row[0], row[1] * 1.0)
)
check how well or how bad our model performed:
import pyspark.mllib.evaluation as ev
LR_evaluation = ev.BinaryClassificationMetrics(LR_results)
print('Area under PR: {0:.2f}'.format(LR_evaluation.areaUnderPR))
print(
 'Area under ROC: {0:.2f}'.format(LR_evaluation.areaUnderROC)
)
LR_evaluation.unpersist()
Here's what we got:
Area under PR: 0.85
Area under ROC: 0.63

```

- Selecting only the most predictable features

```

selector = ft.ChiSqSelector(4).fit(births_train)
topFeatures_train = (
 births_train.map(lambda row: row.label).zip(
 selector.transform(
 births_train.map(lambda row: row.features)
)
)
).map(lambda row: reg.LabeledPoint(row[0], row[1]))
topFeatures_test = (
 births_test.map(lambda row: row.label).zip(
 selector.transform(
 births_test.map(lambda row: row.features)
)
)
).map(lambda row: reg.LabeledPoint(row[0], row[1]))

```
- Predicting infant survival by logistic regression with SGD

```

from pyspark.mllib.tree import RandomForest
RF_model = RandomForest.trainClassifier(
 data=topFeatures_train, numClasses=2,
 categoricalFeaturesInfo={}, numTrees=6,
 featureSubsetStrategy='all', seed=666
)
check how well or how bad our model performed:
RF_results = (
 topFeatures_test.map(lambda row: row.label).zip(
 RF_model.predict(
 topFeatures_test.map(lambda row: row.features)
)
)
)

```

```

)
RF_evaluation = ev.BinaryClassificationMetrics(RF_results)
print('Area under PR: {0:.2f}'.format(RF_evaluation.areaUnderPR))
print(
 'Area under ROC: {0:.2f}'.format(RF_evaluation.areaUnderROC)
)
model_evaluation.unpersist()
Here are the results:
Area under PR: 0.86
Area under ROC: 0.63

```

- Selecting only the most predictable features

```

LR_Model_2 = LogisticRegressionWithLBFGS \
 .train(topFeatures_train, iterations=10)
LR_results_2 = (
 topFeatures_test.map(lambda row: row.label).zip(
 LR_Model_2.predict(
 topFeatures_test.map(lambda row: row.features)
)
)
).map(lambda row: (row[0], row[1] * 1.0))
LR_evaluation_2 = ev.BinaryClassificationMetrics(LR_results_2)
print(
 'Area under PR: {0:.2f}' \
 .format(LR_evaluation_2.areaUnderPR)
)
print(
 'Area under ROC: {0:.2f}' \
 .format(LR_evaluation_2.areaUnderROC)
)
LR_evaluation_2.unpersist()
The results might surprise you:
Area under PR: 0.85
Area under ROC: 0.63

```

### ***The ML Package on DataFrames***

- 3 main abstract classes: a Transformer, an Estimator, and a Pipeline.
- Predicting the chances of infant survival with ML

```

load the data
import pyspark.sql.types as typ
labels = [
 ('INFANT_ALIVE_AT_REPORT', typ.IntegerType()),
 ('BIRTH_PLACE', typ.StringType()),
 ('MOTHER_AGE_YEARS', typ.IntegerType()),
 ('FATHER_COMBINED_AGE', typ.IntegerType()),
 ('CIG_BEFORE', typ.IntegerType()),
 ('CIG_1_TRI', typ.IntegerType()),
 ('CIG_2_TRI', typ.IntegerType()),

```



```

('CIG_3_TRI', typ.IntegerType()),
('MOTHER_HEIGHT_IN', typ.IntegerType()),
('MOTHER_PRE_WEIGHT', typ.IntegerType()),
('MOTHER_DELIVERY_WEIGHT', typ.IntegerType()),
('MOTHER_WEIGHT_GAIN', typ.IntegerType()),
('DIABETES_PRE', typ.IntegerType()),
('DIABETES_GEST', typ.IntegerType()),
('HYP_TENS_PRE', typ.IntegerType()),
('HYP_TENS_GEST', typ.IntegerType()),
('PREV_BIRTH_PRETERM', typ.IntegerType())
]
schema = typ.StructType([
 typ.StructField(e[0], e[1], False) for e in labels
])
births = spark.read.csv(
 'births_transformed.csv.gz', header=True, schema=schema
)
convert to numeric values
import pyspark.ml.feature as ft
births = births \
 .withColumn('BIRTH_PLACE_INT', births['BIRTH_PLACE'] \
 .cast(typ.IntegerType()))
create Transformer:
encoder = ft.OneHotEncoder(
 inputCol='BIRTH_PLACE_INT', outputCol='BIRTH_PLACE_VEC'
)
create a single column with all the features collated together.
featuresCreator = ft.VectorAssembler(
 inputCols=[
 col[0] for col in labels[2:]
] + [encoder.getOutputCol()],
 outputCol='features'
)
Creating an estimator
import pyspark.ml.classification as cl
logistic = cl.LogisticRegression(
 maxIter=10, regParam=0.01, labelCol='INFANT_ALIVE_AT_REPORT'
)
Creating a pipeline
from pyspark.ml import Pipeline
pipeline = Pipeline(stages=[
 encoder, featuresCreator, logistic
])
split dataset into training, validation, and testing datasets.
train, test, val = births.randomSplit([0.7, 0.2, 0.1], seed=666)
Fitting the model
model = pipeline.fit(births_train)
estimation
test_model = model.transform(births_test)

```

```

test_model.take(1)
Evaluating the performance of the model
import pyspark.ml.evaluation as ev
evaluator = ev.BinaryClassificationEvaluator(
 rawPredictionCol='probability',
 labelCol='INFANT_ALIVE_AT_REPORT'
)
print(evaluator.evaluate(
 test_model, {evaluator.metricName: 'areaUnderROC'}
))
print(evaluator.evaluate(
 test_model, {evaluator.metricName: 'areaUnderPR'}
))
save the Pipeline definition for later use
pipelinePath = './infant_oneHotEncoder_Logistic_Pipeline'
pipeline.write().overwrite().save(pipelinePath)
load it up and use it straight away to .fit(...) and predict:
loadedPipeline = Pipeline.load(pipelinePath)
loadedPipeline.fit(births_train).transform(births_test).take(1)
Saving the model
from pyspark.ml import PipelineModel
modelPath = './infant_oneHotEncoder_Logistic_PipelineModel'
model.write().overwrite().save(modelPath)
load it up and use it straight away to predict:
loadedPipelineModel = PipelineModel.load(modelPath)
test_reloadedModel = loadedPipelineModel.transform(births_test)

```

- **Parameter hyper-tuning**
  - **Grid search**

```

import pyspark.ml.tuning as tune
specify the list of parameters we want to loop through:
logistic = cl.LogisticRegression(
 labelCol='INFANT_ALIVE_AT_REPORT'
)
grid = tune.ParamGridBuilder().addGrid(
 logistic.maxIter, [2, 10, 50]
).addGrid(
 logistic.regParam, [0.01, 0.05, 0.3]
).build()
comparing the models:
evaluator = ev.BinaryClassificationEvaluator(
 rawPredictionCol='probability',
 labelCol='INFANT_ALIVE_AT_REPORT'
)
validation
cv = tune.CrossValidator(
 estimator=logistic, estimatorParamMaps=grid,
 evaluator=evaluator
)

```

```

create a purely transforming Pipeline:
pipeline = Pipeline(stages=[encoder ,featuresCreator])
data_transformer = pipeline.fit(births_train)
find the optimal combination of parameters for our model:
cvModel = cv.fit(data_transformer.transform(births_train))
check if it performed better than previous model:
data_train = data_transformer.transform(births_test)
results = cvModel.transform(data_train)
print(evaluator.evaluate(
 results, {evaluator.metricName: 'areaUnderROC'}
))
print(evaluator.evaluate(
 results, {evaluator.metricName: 'areaUnderPR'}
))
What parameters does the best model have?
results = [
 (
 [
 {key.name: paramValue} for key, paramValue in zip(
 params.keys(), params.values()
)
], metric
)
 for params, metric in zip(
 cvModel.getEstimatorParamMaps(), cvModel.avgMetrics
)
]
sorted(results, key=lambda el: el[1], reverse=True)[0]
• Train-validation splitting
selector = ft.ChiSqSelector(
 numTopFeatures=5,
 featuresCol=featuresCreator.getOutputCol(),
 outputCol='selectedFeatures',
 labelCol='INFANT_ALIVE_AT_REPORT'
)
logistic = cl.LogisticRegression(
 labelCol='INFANT_ALIVE_AT_REPORT',
 featuresCol='selectedFeatures'
)
pipeline = Pipeline(stages=[encoder, featuresCreator, selector])
data_transformer = pipeline.fit(births_train)
tvs = tune.TrainValidationSplit(
 estimator=logistic,
 estimatorParamMaps=grid,
 evaluator=evaluator
)
tvsModel = tvs.fit(
 data_transformer.transform(births_train)

```

```


)
data_train = data_transformer.transform(births_test)
results = tvsModel.transform(data_train)
print(evaluator.evaluate(
 results, {evaluator.metricName: 'areaUnderROC'}
))
print(evaluator.evaluate(
 results, {evaluator.metricName: 'areaUnderPR'}
))
• NLP - related feature extractors
text dataset
text_data = spark.createDataFrame(
 [
 ['Machine learning can be applied to a wide variety of
 data types, such as vectors, text, images, and structured
 data. This API adopts the DataFrame from Spark SQL in order
 to support a variety of data types.'],
 ...
 ['Columns in a DataFrame are named. The code examples
 below use names such as "text," "features," and "label."']
],
 ['input']
)
tokenization
tokenizer = ft.RegexTokenizer(
 inputCol='input', outputCol='input_arr', pattern='\s+|[,.\"]'
)
remove stopwords
stopwords = ft.StopWordsRemover(
 inputCol=tokenizer.getOutputCol(), outputCol='input_stop'
)
build NGram model and the Pipeline:
ngram = ft.NGram(
 n=2, inputCol=stopwords.getOutputCol(), outputCol="nGrams"
)
pipeline = Pipeline(stages=[tokenizer, stopwords, ngram])
data_ngram = pipeline.fit(text_data).transform(text_data)
data_ngram.select('nGrams').take(1)
• Discretizing highly non-linear continuous variables into discrete buckets
import numpy as np
x = np.arange(0, 100)
x = x / 100.0 * np.pi * 4
y = x * np.sin(x / 1.764) + 20.1234
create a DataFrame
schema = typ.StructType([
 typ.StructField('continuous_var', typ.DoubleType(), False)
])
data = spark.createDataFrame(

```

```

[[float(e),] for e in y], schema=schema
)
split continuous variable
discretizer = ft.QuantileDiscretizer(
 numBuckets=5, inputCol='continuous_var',
 outputCol='discretized'
)
fit the data
data_discretized = discretizer.fit(data).transform(data)
Standardizing continuous variables
vectorizer = ft.VectorAssembler(
 inputCols=['continuous_var'], outputCol= 'continuous_vec'
)
normalization
normalizer = ft.StandardScaler(
 inputCol=vectorizer.getOutputCol(), outputCol='normalized',
 withMean=True, withStd=True
)
pipeline = Pipeline(stages=[vectorizer, normalizer])
data_standardized = pipeline.fit(data).transform(data)
Here's what the transformed data would look like:

```



### GraphFrames

```

Preparing your flights dataset
Airline On-Time Performance and Causes of Flight Delays:
contains scheduled and actual departure and arrival times, and
delay causes as reported by US air carriers.
Open Flights: Airports and airline data:
including the IATA code, airport name, and airport location.
Set File Paths
tripdelaysFilePath =

```

```

"/databricksdatasets/flights/departuredelays.csv"
airportsnaFilePath =
"/databricks-datasets/flights/airportcodes-na.txt"
Obtain airports dataset
Note, this dataset is tab-delimited with a header
airportsna = spark.read.csv(
 airportsnaFilePath, header='true', inferSchema='true',
 sep='\t'
)
airportsna.createOrReplaceTempView("airports_na")
Obtain departure Delays data
Note, this dataset is comma-delimited with a header
departureDelays = spark.read.csv(
 tripdelaysFilePath, header='true'
)
departureDelays.createOrReplaceTempView("departureDelays")
departureDelays.cache()
Available IATA codes from the departedelays sample dataset
tripIATA = spark.sql(
 "select distinct iata from (
 select distinct origin as iata from departureDelays union
 all select distinct destination as iata from
 departureDelays
)"
)
tripIATA.createOrReplaceTempView("tripIATA")
Only include airports with at least one trip from the
`departureDelays` dataset
airports = spark.sql(
 "select f.IATA, f.City, f.State, f.Country from airports_na f
 join tripIATA t on t.IATA = f.IATA"
)
airports.createOrReplaceTempView("airports")
airports.cache()
Build `departureDelays_geo` DataFrame
Obtain key attributes such as Date of flight, delays, distance,
and airport information (Origin, Destination)
departureDelays_geo = spark.sql(
 "select cast(f.date as int) as tripid,
 cast(concat(concat(concat(concat(concat(concat('2014-',
 concat(concat(substr(cast(f.date as string), 1, 2), '-')),
 substr(cast(f.date as string), 3, 2)), '')), substr(cast(f.date as string), 5, 2)), ':'),
 substr(cast(f.date as string), 7, 2)), ':00') as timestamp) as `localdate`, cast(f.delay as int),
 cast(f.distance as int), f.origin as src, f.destination as dst,
 o.city as city_src, d.city as city_dst, o.state as state_src,
 d.state as state_dst from departedelays f join airports o on
 o.iata = f.origin join airports d on d.iata = f.destination"
)

```

```

Create Temporary View and cache
departureDelays_geo.createOrReplaceTempView(
 "departureDelays_geo"
)
departureDelays_geo.cache()
Review the top 10 rows of the `departureDelays_geo` DataFrame
departureDelays_geo.show(10)
Building the graph
Note, ensure you have already installed GraphFrames spark-package
from pyspark.sql.functions import *
from graphframes import *
Create Vertices (airports) and Edges (flights)
tripVertices = airports.withColumnRenamed("IATA", "id").distinct()
tripEdges = departureDelays_geo.select(
 "tripid", "delay", "src", "dst", "city_dst", "state_dst"
)
Cache Vertices and Edges
tripEdges.cache()
tripVertices.cache()
display(tripEdges)
create a GraphFrame using the GraphFrame command:
tripGraph = GraphFrame(tripVertices, tripEdges)
Executing simple queries
print "Airports: %d" % tripGraph.vertices.count()
print "Trips: %d" % tripGraph.edges.count()
Determining the longest delay in this dataset
tripGraph.edges.groupBy().max("delay")
Determining the number of delayed versus on-time/early flights
print "On-time / Early Flights: %d" % tripGraph.edges.filter("delay
<= 0").count()
print "Delayed Flights: %d" % tripGraph.edges.filter("delay >
0").count()
What flights departing Seattle likely to have significant delays?
tripGraph.edges\
 .filter("src = 'SEA' and delay > 0")\
 .groupBy("src", "dst")\
 .avg("delay")\
 .sort(desc("avg(delay)"))\
 .show(5)
What states have significant delays departing from Seattle?
States with the longest cumulative delays (with individual
delays > 100 minutes) (origin: Seattle)
display(tripGraph.edges.filter("src = 'SEA' and delay > 100"))
Understanding vertex degrees
display(tripGraph.degrees.sort(desc("degree")).limit(20))
display(tripGraph.inDegrees.sort(desc("inDegree")).limit(20))
display(tripGraph.outDegrees.sort(desc("outDegree")).limit(20))
Determining the top transfer airports by the ratio
Calculate the inDeg (flights into the airport) and

```

```

outDeg (flights leaving the airport)
inDeg = tripGraph.inDegrees
outDeg = tripGraph.outDegrees
Calculate the degreeRatio (inDeg/outDeg)
degreeRatio = inDeg.join(outDeg, inDeg.id == outDeg.id) \
 .drop(outDeg.id) \
 .selectExpr(
 "id", "double(inDegree)/double(outDegree) as degreeRatio"
) \
 .cache()
Join back to the 'airports' DataFrame
(instead of registering temp table as above)
transferAirports = degreeRatio.join(
 airports, degreeRatio.id == airports.IATA
).selectExpr("id", "city", "degreeRatio") \
 .filter("degreeRatio between 0.9 and 1.1")
List out the top 10 transfer city airports
display(transferAirports.orderBy("degreeRatio").limit(10))
Understanding motifs
Generate motifs
motifs = tripGraphPrime.find("(a)-[ab]->(b); (b)-[bc]->(c)")\
 .filter("(b.id = 'SFO') and (ab.delay > 500 or bc.delay >
500) and bc.tripid > ab.tripid and bc.tripid < ab.tripid +
10000")
Display motifs
display(motifs)
Determining airport ranking using PageRank
Determining Airport ranking of importance using 'pageRank'
ranks = tripGraph.pageRank(resetProbability=0.15, maxIter=5)
Display the pageRank output
display(
 ranks.vertices.orderBy(ranks.vertices.pagerank.desc()) \
 .limit(20)
)
Determining the most popular non-stop flights
Determine the most popular non-stop flights
import pyspark.sql.functions as func
topTrips = tripGraph \
 .edges \
 .groupBy("src", "dst") \
 .agg(func.count("delay").alias("trips"))
Show the top 20 most popular flights (single city hops)
display(topTrips.orderBy(topTrips.trips.desc()).limit(20))
Using Breadth-First Search
Obtain list of direct flights between SEA and SFO
filteredPaths = tripGraph.bfs(
 fromExpr = "id = 'SEA'", toExpr = "id = 'SFO'",
 maxPathLength = 1
)

```

```

display list of direct flights
display(filteredPaths)
Obtain list of direct flights between SFO and BUF
filteredPaths = tripGraph.bfs(
 fromExpr = "id = 'SFO'", toExpr = "id = 'BUF'",
 maxPathLength = 1
)
display list of direct flights
display(filteredPaths)
display list of one-stop flights between SFO and BUF
filteredPaths = tripGraph.bfs(
 fromExpr = "id = 'SFO'", toExpr = "id = 'BUF'",
 maxPathLength = 2
)
display list of flights
display(filteredPaths)
Display most popular layover cities by descending count
display(
 filteredPaths.groupBy("v1.id", "v1.City") \
 .count().orderBy(desc("count")).limit(10)
)
Visualizing flights using D3
%scala
// On-time and Early Arrivals
import d3a._
graphs.force(
 height = 800, width = 1200,
 clicks = sql("""
 select src, dst as dest, count(1) as count from
 departureDelays_geo where delay <= 0 group by src, dst
 """).as[Edge]
)

```

### TensorFrames

- Matrix multiplication using constants

```

Import TensorFlow
import tensorflow as tf
Setup the matrix
c1: 1x3 matrix
c2: 3x1 matrix
c1 = tf.constant([[3., 2., 1.]])
c2 = tf.constant([[-1.], [2.], [1.]])
tensors in the form of numpy ndarray or
tensorflow::Tensor interfaces in C/C++
m3: matrix multiplication (m1 x m3)
mp = tf.matmul(c1, c2)
Launch the default graph
s = tf.Session()

```

```

run: Execute the ops in graph
r = s.run(mp)
print(r)
Close the Session when completed
s.close()

```

- Matrix multiplication using placeholders for tensors of different sizes/shape

```

Setup placeholder for your model
t1: placeholder tensor
t2: placeholder tensor
t1 = tf.placeholder(tf.float32)
t2 = tf.placeholder(tf.float32)
t3: matrix multiplication (m1 x m3)
tp = tf.matmul(t1, t2)
Running the model
Define input matrices
m1 = [[3., 2., 1.]]
m2 = [[-1.], [2.], [1.]]
Execute the graph within a session
with tf.Session() as s:
 print(s.run([tp], feed_dict={t1:m1, t2:m2}))
setup input matrices
m1 = [[3., 2., 1., 0.]]
m2 = [[-5.], [-4.], [-3.], [-2.]]
Execute the graph within a session
with tf.Session() as s:
 print(s.run([tp], feed_dict={t1:m1, t2:m2}))

```

- TensorFrames – quick start

```

The version we're using in this notebook
$SPARK_HOME/bin/pyspark --packages \
tjhunter:tensorframes:0.2.2-s_2.10
Or use the latest version
$SPARK_HOME/bin/pyspark --packages \
databricks:tensorframes:0.2.3-s_2.10

```

- Using TensorFlow to add a constant to an existing column

```

Import TensorFlow, TensorFrames, and Row
import tensorflow as tf
import tensorframes as tfs
from pyspark.sql import Row
Create RDD of floats and convert into DataFrame `df`
rdd = [Row(x=float(x)) for x in range(10)]
df = sqlContext.createDataFrame(rdd)
df.show()
Executing the Tensor graph
Run TensorFlow program executes:
The 'op' performs the addition (i.e. 'x' + '3')
Place the data back into a DataFrame
with tf.Graph().as_default() as g:
 # The placeholder that corresponds to column 'x'. The shape of

```

```

the placeholder is automatically inferred from the DataFrame.
x = tfs.block(df, "x")
The output that adds 3 to x
z = tf.add(x, 3, name='z')
The resulting `df2` DataFrame
df2 = tfs.map_blocks(z, df)
Note that 'z' is the tensor output from the 'tf.add' operation
print z
Output
Tensor("z:0", shape=(?,), dtype=float64)

```

- **Blockwise reducing operations example**

```

Build a DataFrame of vectors
data = [Row(y=[float(y), float(-y)]) for y in range(10)]
df = sqlContext.createDataFrame(data)
df.show()
Analysing the DataFrame
Print the information gathered by TensorFlow
tfs.print_schema(df)
Output
root
|-- y: array (nullable = true) double[?,?]
Because the dataframe contains vectors, we need to analyze it
first to find the dimensions of the vectors.
df2 = tfs.analyze(df)
The information gathered by TF can be printed to check the content
tfs.print_schema(df2)
Output
root
|-- y: array (nullable = true) double[?,2]
Computing elementwise sum (.reduce_sum) and min (.reduce_min)
Note: First, let's make a copy of the 'y' column.
This is an inexpensive operation in Spark 2.0+
df3 = df2.select(df2.y, df2.y.alias("z"))
Execute the Tensor Graph
with tf.Graph().as_default() as g:
 # The placeholders.
 # Note the special name that end with '_input':
 y_input = tfs.block(df3, 'y', tf_name="y_input")
 z_input = tfs.block(df3, 'z', tf_name="z_input")
 # Perform elementwise sum and minimum
 y = tf.reduce_sum(y_input, [0], name='y')
 z = tf.reduce_min(z_input, [0], name='z')
The resulting dataframe
(data_sum, data_min) = tfs.reduce_blocks([y, z], df3)
The final results are numpy arrays:
print "Elementwise sum: %s and min: %s " % (data_sum, data_min)
Output
Elementwise sum: [45. -45.] and minimum: [0. -9.]

```

### Polyglot Persistence with Blaze

- **Working with NumPy arrays**

```

import numpy as np
simpleArray = np.array([[1,2,3], [4,5,6]])
simpleData_np = bl.Data(simpleArray)
simpleData_np.peak()
simpleData_np[0]
transpose your DataShape: but the name of the column becomes None
simpleData_np.T[0]
named columns => calling the column by its name
simpleData_np = bl.Data(simpleArray, fields=['a', 'b', 'c'])
simpleData_np['b']

```
- **Working with pandas' DataFrame**

```

import pandas as pd
simpleDf = pd.DataFrame([[1,2,3], [4,5,6]], columns=['a','b','c'])
transform it into a DataShape:
simpleData_df = bl.Data(simpleDf)
retrieve data
simpleData_df['a']

```
- **Working with files**

```

import odo
traffic = bl.Data('../Data/TrafficViolations.csv')
print(traffic.fields)
traffic_gz = bl.Data('../Data/TrafficViolations.csv.gz')
produces the same results, it takes more time to retrieve from
the archived file
traffic.head(2)
traffic_gz.head(2)
saves the data into a GZip archive
for year in traffic.Stop_year.distinct().sort():
 odo.odo(
 traffic[traffic.Stop_year == year],
 '../Data/Years/TrafficViolations_{0}.csv.gz'.format(year)
)
read from multiple files using the asterisk character *:
traffic_multiple = bl.Data(
 '../Data/Years/TrafficViolations_*.csv.gz'
)
traffic_multiple.head(2)
you can read data from JSON, Excel files, HDFS, or bcolz files.

```
- **Interacting with relational databases**

```

traffic_psql = bl.Data(
 'postgres://{{0}}:{{1}}@localhost:5432/drabast::traffic\'
 .format('<your_username>', '<your_password>')
)
traffic_2016 = traffic_psql[traffic_psql['Year'] == 2016]
Drop commands

```



```

#odo.drop('sqlite:///traffic_local.sqlite::traffic2016')
#odo.drop('postgresql://{0}:{1}@localhost:5432/drabast::traffic'\
.format('<your_username>', '<your_password>'))
Save to SQLite
odo.odo(
 traffic_2016, 'sqlite:///traffic_local.sqlite::traffic2016'
)
Save to PostgreSQL
odo.odo(
 traffic_2016,
 'postgresql://{0}:{1}@localhost:5432/drabast::traffic'\
 .format('<your_username>', '<your_password>')
)
read from SQLite
traffic_sqlt = bl.Data(
 'sqlite:///traffic_local.sqlite::traffic2016'
)

```

- Interacting with the MongoDB database

```

read from MongoDB
traffic_mongo = bl.Data(
 'mongodb://localhost:27017/packt::traffic'
)

```
- Data operations

```

Accessing a single column
traffic.Year.head(2)
selection of more than one column at a time:
(traffic[['Location', 'Year', 'Accident', 'Fatal',
'Alcohol']].head(2))
the equivalent SQL code would be:
SELECT *
FROM traffic
LIMIT 2z

```
- Symbolic transformations: Blaze can operate symbolically using the .symbol() method; the first argument specifies a symbolic name of the transformation, and the second argument specifies the schema

```

schema_example = bl.symbol(
 'schema_examp1', '{id: int, name: string}'
)
reuse the schema by using traffic.dshape
traffic_s = bl.symbol('traffic', traffic.dshape)
traffic_2013 = traffic_s[traffic_s['Stop_year'] == 2013][
 ['Stop_year', 'Arrest_Type', 'Color', 'Charge']]
traffic_pd = pd.read_csv('../Data/TrafficViolations.csv')
perform the computation using the .compute() method:
the first argument specifies the transformation object and
the second parameter is the data
bl.compute(traffic_2013, traffic_pd).head(2)

```

```

You can also pass a list of lists or a list of NumPy arrays
using the .values attribute
bl.compute(traffic_2013, traffic_pd.values)[0:2]

```

- Operations on columns

```

mathematical operations to be done on numeric columns.
traffic['Stop_year'].distinct().sort()
An equivalent syntax for pandas would be as follows:
traffic['Stop_year'].unique().sort()
For SQL, use the following code:
SELECT DISTINCT Stop_year
FROM traffic
mathematical transformations/arithmetic to the columns.
traffic['Stop_year'].head(2) - 2000
For SQL, the equivalent would be:
SELECT Stop_year - 2000 AS Stop_year
FROM traffic
more complex mathematical operations (for example, log or pow)
bl.log(traffic['Stop_year']).head(2)

```
- Reducing data

```

reduction methods, such as .mean(), .std, or .max()
traffic['Stop_year'].max()
for SQL the same could be done with the following code:
SELECT MAX(Stop_year) AS Stop_year_max
FROM traffic
the .transform() method
traffic = bl.transform(
 traffic, Age_of_car = traffic.Stop_year - traffic.Year
)
traffic.head(2)
An equivalent operation in pandas
traffic['Age_of_car'] = traffic.apply(
 lambda row: row.Stop_year - row.Year, axis = 1
)
For SQL you can use the following code:
SELECT *, Stop_year - Year AS Age_of_car
FROM traffic
perform a group by operation using the .by() operation:
bl.by(
 traffic['Fatal'],
 Fatal_AvgAge=traffic.Age_of_car.mean(),
 Fatal_Count =traffic.Age_of_car.count()
)
For pandas, an equivalent would be as follows:
traffic.groupby('Fatal')['Age_of_car']\
 .agg({
 'Fatal_AvgAge': np.mean, 'Fatal_Count': np.count_nonzero
 })
For SQL, it would be as follows:

```

```

SELECT Fatal
 , AVG(Age_of_car) AS Fatal_AvgAge
 , COUNT(Age_of_car) AS Fatal_Count
FROM traffic
GROUP BY Fatal

```

- Joins
  - # Joining two DataShapes using .join()
 

```

violation = traffic[[
 'Stop_month', 'Stop_day', 'Stop_year',
 'Stop_hr', 'Stop_min', 'Stop_sec', 'Violation_Type'
]]
belts = traffic[[
 'Stop_month', 'Stop_day', 'Stop_year',
 'Stop_hr', 'Stop_min', 'Stop_sec', 'Belts'
]]
violation_belts = bl.join(
 violation, belts, [
 'Stop_month', 'Stop_day', 'Stop_year',
 'Stop_hr', 'Stop_min', 'Stop_sec'
]
)
bl.by(
 violation_belts[['Violation_Type', 'Belts']],
 Violation_count=violation_belts.Belts.count()
).sort('Violation_count', ascending=False)
The same could be achieved in pandas with the following code:
violation.merge(
 belts,
 on=[
 'Stop_month', 'Stop_day', 'Stop_year',
 'Stop_hr', 'Stop_min', 'Stop_sec'
]
).groupby(['Violation_type', 'Belts']).agg({
 'Violation_count': np.count_nonzero
}) \
.sort('Violation_count', ascending=False)
With SQL, you would use the following snippet:
SELECT innerQuery.* FROM (
 SELECT a.Violation_type, b.Belts, COUNT() AS Violation_count
 FROM violation AS a
 INNER JOIN belts AS b
 ON a.Stop_month = b.Stop_month
 AND a.Stop_day = b.Stop_day
 AND a.Stop_year = b.Stop_year
 AND a.Stop_hr = b.Stop_hr
 AND a.Stop_min = b.Stop_min
 AND a.Stop_sec = b.Stop_sec
 GROUP BY Violation_type, Belts

```

```

) AS innerQuery
ORDER BY Violation_count DESC

```

### Structured Streaming

- Spark Streaming word count application using DStreams and Unix/Linux nc command (read and write data across network connection).
  - # Create a local SparkContext and Streaming Contexts
 

```

from pyspark import SparkContext
from pyspark.streaming import StreamingContext

```
  - # Create sc with two working threads
 

```

sc = SparkContext("local[2]", "NetworkWordCount")

```
  - # Create local StreamingContext with batch interval of 1 second
 

```

ssc = StreamingContext(sc, 1)

```
  - # Create DStream that connects to localhost:9999
 

```

lines = ssc.socketTextStream("localhost", 9999)

```
  - # Split lines into words
 

```

words = lines.flatMap(lambda line: line.split(" "))

```
  - # Count each word in each batch
 

```

pairs = words.map(lambda word: (word, 1))

```
  - wordCounts = pairs.reduceByKey(lambda x, y: x + y)
 

```

wordCounts = pairs.reduceByKey(lambda x, y: x + y)

```
  - # Print the first ten elements of each RDD in this DStream
 

```

wordCounts.pprint()

```
  - # Start the computation
 

```

ssc.start()

```
  - # Wait for the computation to terminate with <Ctrl><C>
 

```

ssc.awaitTermination()

```

To start the nc command, from one of your terminals:

```

$ nc -lk 9999

```
- Global aggregations: calculating a stateful aggregation beyond batch interval with windowing
  - # inserting any new chunks of data led to slow streaming
  - # performance with the ever-increasing scheduling delays.

```

sqlContext.sql(
 "insert into meetup_stream select * from meetup_stream_json"
)

```

  - # creating global aggregations via UpdateStateByKey (Spark 1.5)
  - # the performance is proportional to the size of the state.
  - # Create a local SparkContext and Streaming Contexts
 

```

from pyspark import SparkContext
from pyspark.streaming import StreamingContext

```
  - # Create sc with two working threads
 

```

sc = SparkContext("local[2]", "StatefulNetworkWordCount")

```
  - # Create local StreamingContext with batch interval of 1 sec
 

```

ssc = StreamingContext(sc, 1)

```
  - # Create checkpoint for local StreamingContext: ensure that Spark
  - # Streaming is fault tolerant
 

```

ssc.checkpoint("checkpoint")

```

```
Define updateFunc via UpdateStateByKey: sum of (key, value) pairs
def updateFunc(new_values, last_sum):
 return sum(new_values) + (last_sum or 0)
Create DStream that connects to localhost:9999
lines = ssc.socketTextStream("localhost", 9999)
Calculate running counts
running_counts = lines.flatMap(lambda line: line.split(" "))\
 .map(lambda word: (word, 1)).updateStateByKey(updateFunc)
Print the first ten elements of each RDD generated in this
stateful DStream to the console
running_counts.pprint()
Start the computation
ssc.start()
Wait for the computation to terminate
ssc.awaitTermination()
creating global aggregations via mapWithState (Spark 1.6)
The performance is proportional to the size of the batch.
```

- Structured Streaming (Spark 2.0)

- batch aggregation
 

```
reads a data stream from S3 and saves it to a MySQL database:
logs = spark.read.json('s3://logs')
logs.groupBy(logs.UserId).agg(sum(logs.Duration)) \
 .write.jdbc('jdbc:mysql://...')
continuous aggregation:
logs = spark.readStream.json('s3://logs').load()
sq = logs.groupBy(logs.UserId).agg(sum(logs.Duration)) \
 .writeStream.format('json').start()
Will return true if the `sq` stream is active
sq.isActive
Will terminate the `sq` stream
sq.stop()
```
- DataFrames code
 

```
Import the necessary classes and create a local SparkSession
from pyspark.sql import SparkSession
from pyspark.sql.functions import explode
from pyspark.sql.functions import split
spark = SparkSession.builder \
 .appName("StructuredNetworkWordCount").getOrCreate()
Create DataFrame representing the stream of input lines
from connection to localhost:9999
lines = spark\
 .readStream\
 .format('socket')\
 .option('host', 'localhost')\
 .option('port', 9999)\
 .load()
Split the lines into words
words = lines.select(
```

```
explode(
 split(lines.value, ' ')
).alias('word')
)
Generate running word count
wordCounts = words.groupBy('word').count()
• output this data to the console
Start running the query that prints the
running counts to the console
query = wordCounts\
 .writeStream\
 .outputMode('complete')\
 .format('console')\
 .start()
Await Spark Streaming termination
query.awaitTermination()
• run nc job in the first terminal:
$ nc -lk 9999
```

### Packaging Spark Applications

- submitting jobs to Spark using the spark-submit script.
 

```
$ spark-submit [options] <python file> [app arguments]
```

--master: Parameter used to set the URL of the master (head) node: Local, local[n], local[\*], Spark standalone cluster spark://host:port, mesos://host:port, and Yarn.

--deploy-mode: client or cluster

--name: Name of your application.

--py-files: Comma-delimited list of .py, .egg or .zip files to include for Python apps.

--files: comma-delimited list of files.

--conf: configure of your app dynamically from the command line parameters, or specified in the conf/spark-defaults.conf file. The syntax is <Spark property>=<value for the property>.

--properties-file: File with a configuration having the same set of properties as the conf/spark-defaults.conf file.

--driver-memory: default is 1,024M.

--executor-memory: default is 1G.

--help:

--verbose:

--version:

--supervise:

--kill:

--status:
- Deploying the app programmatically

```
Configuring your SparkSession and creating SparkSession
from pyspark.sql import SparkSession
spark = SparkSession \
 .builder \
 .appName('CalculatingGeoDistances') \
 .getOrCreate()
print('Session created')
```

- Modularizing code

- Structure of the Python package:

```
additionalCode/
 setup.py
 utilities/
 __init__.py
 base.py
 converters/
 __init__.py
 distance.py
 geoCalc.py
```

- The setup.py file in our case looks as follows:

```
from setuptools import setup
setup(
 name='PySparkUtilities',
 version='0.1dev',
 packages=['utilities', 'utilities/converters'],
 license='''
 Creative Commons
 Attribution-Noncommercial-Share Alike license''',
 long_description='''
 An example of how to package code for PySpark'''
)
```

- The \_\_init\_\_.py file in the utilities folder has the following code:

```
from .geoCalc import geoCalc
__all__ = ['geoCalc', 'converters']
```

- Calculating the distance in miles between two points (latitude and longitude) on a map (Cartesian coordinates) using the Haversine formula with calculateDistance() in the geoCalc.py file.

- Converting distance units: any class implemented as a converter should expose the same interface and implement the convert() method. E.g.,

```
from abc import ABCMeta, abstractmethod
class BaseConverter(metaclass=ABCMeta):
 @staticmethod
 @abstractmethod
 def convert(f, t):
 raise NotImplementedError
```

- Building an egg

```
$ python setup.py bdist_egg
```

- User defined functions in Spark

```
import utilities.geoCalc as geo
from utilities.converters import metricImperial
getDistance = func.udf(
 lambda lat1, long1, lat2, long2:
 geo.calculateDistance(
 (lat1, long1), (lat2, long2)
)
)
convertMiles = func.udf(lambda m:
 metricImperial.convert(str(m) + ' mile', 'km'))
```

- calculate the distance and convert it to miles:

```
Using the .withColumn() method we create additional columns.
uber = uber.withColumn(
 'miles',
 getDistance(
 func.col('pickup_latitude'),
 func.col('pickup_longitude'),
 func.col('dropoff_latitude'),
 func.col('dropoff_longitude')
)
)
uber = uber.withColumn(
 'kilometers',
 convertMiles(func.col('miles'))
)
```

- Submitting a job

```
$./launch_spark_submit.sh \
--master local[4] \
--py-files calculatingGeoDistance.py \
additionalCode/dist/PySparkUtilities-0.1.dev0-py3.5.egg
```

- configured Spark instance to run Jupyter and automated it with the launch\_spark\_submit.sh script:

```
#!/bin/bash
unset PYSPARK_DRIVER_PYTHON
spark-submit $*
export PYSPARK_DRIVER_PYTHON=jupyter
```

- Monitoring execution: you can switch between the Jobs view or the Stages view to track all the stages that are executed.